

# CSCI 509 - Operating Systems, Fall 2018

## Assignment 2: Threads and Synchronization

Due Date: October 12, 2018

Points: 200

### 1 Overview and Goal

In this assignment you will use threads and gain experience writing programs involving concurrency control. You will begin by studying the thread package, which implements multi-threading. You will make modifications and additions to the existing code. Then you will use the threads package to solve several concurrency problems.

In addition, you will gain familiarity programming in the KPL language.

If you have difficulties with this assignment or want to go into more depth, take a look at the document called "The Thread Scheduler and Concurrency Control Primitives":

<https://facultyweb.cs.wvu.edu/~phil/classes/blitz/BlitzDoc/ThreadScheduler.htm>

### 2 Download the Files

Start by creating a new branch in your gitlab project. Checkout the new branch and then create a directory named *a2*

Then download all the files from:

<http://facultyweb.cs.wvu.edu/~phil/classes/f18/509/a2>

into your directory. You should get the following files:

```
makefile
DISK
System.h
System.k
Runtime.s
Switch.s
List.h
List.k
Thread.h
Thread.k
Main.h
Main.k
Synch.h
Synch.k
Game.h
Game.k
Barber.h
Barber.k
SleepingBarberProblem.pdf
```

In this assignment, you will modify the following files:

```
Main.k
Synch.h
Synch.k
Game.h
Game.k
Barber.h
Barber.k
```

After getting the files, you should be able to compile all the code (as is) with the UNIX make command. The program executable we are building will be called "os". You can execute the first program by typing:

```
% make
% blitz -g os
```

It also builds "game" that can be executed in a similar way. In the course of your experimentation, you may modify other files besides Synch.h, Synch.k, Main.k, and Game.\* but the code you are required to write and turn in doesn't require any changes to the other files other than Synch.h, Synch.k, Main.k, Game.\* and Barber.\*. For example, you may wish to uncomment some of the print statements, to see what happens. *However, your final versions of the files you change must work with the other files, exactly as they are distributed.*

Be sure you copy all files. Even though there are similarities with some of the files used for the previous assignment, there may be some subtle but important differences.

### 3 Study the Existing Code

The code provided in this assignment provides the ability to create and run multiple threads, and to control concurrency through several synchronization methods.

Start by looking over the System package. Focus on the material toward the beginning of file System.k, namely the functions:

```
print
printInt
printHex
printChar
printBool
nl
MemoryEqual
StrEqual
StrCopy
StrCmp
Min
```

```
Max
printIntVar
printHexVar
printBoolVar
printCharVar
printPtr
```

Get familiar with the printing functions; you'll be calling them a lot. Some are implemented in assembly code and some are implemented in KPL in the System package.

The following functions are used to implement the heap in KPL:

```
KPLSystemInitialize
KPLMemoryAlloc
KPLMemoryFree
```

Objects can be allocated on the heap and freed with the alloc and free statements. The HEAP implementation is very rudimentary in this implementation. In your kernel, you may allocate objects during start-up but after that, **YOU MUST NOT ALLOCATE OBJECTS ON THE HEAP!** Why? Because the heap might fill up and then what is a kernel supposed to do? Crash!

In this assignment, you should not allocate anything on the heap.

The following functions can be ignored since they concern aspects of the KPL language that we will not be using:

```
KPLUncaughtThrow UncaughtThrowError KPLIsKindOf KPLSystemError
```

The Runtime.s file contains a number of routines coded in assembly language. It contains the program entry point and the interrupt vector in low memory. Take a look at it. Follow what happens when program execution begins at location 0x00000000 (the label "\_entry"). The code labeled "\_mainEntry" is included in code the compiler produces. The "\_mainEntry" code will call the main function, which appears in the file Main.k.

In Runtime.s, follow what happens when a timer interrupt occurs. It makes an "up-call" to a routine called \_P\_Thread\_TimerInterruptHandler. This name refers to "a function called TimerInterruptHandler in a package called Thread." (\_P\_Thread\_TimerInterruptHandler is the name the compiler will give this function.)

All the code in this assignment assumes that no other interrupt types (such as a DiskInterrupt) occur. In Runtime.s, follow what would happen if another sort of interrupt should ever occur.

The KPL language will check for lots of error conditions, such as the use of a null pointer. Try changing the program to make this error. Follow in Runtime.s what happens when this occurs.

Next take a look at the List package. Read the header file carefully. This package provides code that implements a linked list. We'll use linked lists in this assignment. For example, the threads that are ready to run (and waiting for time on the CPU) will be kept in a linked list called the "ready list". Threads that become BLOCKED will sit on other linked lists. Also look over the List.k code file.

The key class in this assignment is named `Thread`, and it is located in the `Thread` package along with other stuff (see `Thread.h`, `Thread.k`). For each thread, there will be a single `Thread` object. `Thread` is a subclass of `Listable`, which means that each `Thread` object contains a next pointer and can be added to a linked list.

The `Thread` package is central and you should study this code thoroughly. This package contains one class (called `Thread`) and several functions related to thread scheduling and time-slicing:

```
InitializeScheduler ()
IdleFunction (arg: int)
Run (nextThread: ptr to Thread)
PrintReadyList ()
ThreadStartMain ()
ThreadFinish ()
FatalError (errorMessage: ptr to array of char)
SetInterruptsTo (newStatus: int) returns int
TimerInterruptHandler ()
```

`FatalError` is the simplest function. We will call `FatalError` whenever we wish to print an error message and abort the program. Typically, we'll call `FatalError` after making some check and finding that things are not as expected. `FatalError` will print the name of the thread invoking it, print the message, and then shut down. It will throw us into the BLITZ emulator command line mode. Normally, the next thing to do might be to type the "st" command (short for "stack"), to see which functions and methods were active.

(Of course the information printed out by the emulator will pertain to only the thread that invoked `FatalError`. The emulator does not know about threads, and it is pretty much impossible to extract information about other threads by examining bytes in memory.)

The next function to look at is `SetInterruptsTo`, which is used to change the "I" interrupt bit in the CPU. We can use it to disable interrupts with code like this:

```
... = SetInterruptsTo (DISABLED)
and we can use it to enable interrupts:
... = SetInterruptsTo (ENABLED)
```

This function returns the previous status. This is very useful because we often want to DISABLE interrupts (regardless of what they were before) and then later we want to return the interrupt status to whatever it was before. In our kernel, we'll often see code like:

```
var oldIntStat: int ... oldIntStat = SetInterruptsTo (DISABLED) ... oldIntStat = SetInter-
ruptsTo (oldIntStat)
```

Next take a look at the `Thread` class. Here are the fields of `Thread`:

```
name: ptr to array of char
status: int
systemStack: array [SYSTEM_STACK_SIZE] of int
regs: array [13] of int
stackTop: ptr to void
```

```
initialFunction: ptr to function (int)
initialArgument: int
```

Here are the operations (i.e., methods) you can do on a Thread:

```
Init (n: ptr to array of char)
Fork (fun: ptr to function (int), arg: int)
Yield ()
Sleep ()
CheckOverflow ()
Print ()
```

Each thread is in one of the following states: `JUST_CREATED`, `READY`, `RUNNING`, `BLOCKED`, and `UNUSED`, and this is given in the status field. (The `UNUSED` status is given to a Thread after it has terminated. We'll need this in later assignments.)

Each thread has a name. To create a thread, you'll need a Thread variable. First, use `Init` to initialize it, providing a name.

Each thread needs its own stack and space for this stack is placed directly in the Thread object in the field called `systemStack`. Currently, this is an array of 1000 words, which should be enough. (It is conceivable our code could overflow this limit; there is a check to make sure we don't overflow this limited area.)

All threads in this assignment will run in System mode. Therefore the stack is called the "system stack". In later assignments, we'll see that this stack is used only for kernel routines. User programs will have their own stacks in their virtual address spaces, but we are getting ahead of ourselves.

The Thread object also has space to store the state of the CPU, namely the registers. Whenever a thread switch occurs, the registers will be saved in the Thread object. These fields (`regs` and `stackTop`) are used by the assembly code routine named `Switch`.

The Thread object also has space to store a pointer to a function (the `initialFunction` field) and an argument for this function (the `initialArgument` field). This pointer will point to the "main" function of this thread; this is the function that will get executed when this thread begins execution. We are storing a pointer to the function because this is a variable: different threads may execute different functions.

We are also able to supply an initial argument to this thread, through the `initialArgument` field. This argument must be an integer. Often there will be several threads executing the same "main" function. The argument is a handy way to let each thread know what its role should be. For example, we might create 10 threads each using the same "main" function, but passing each thread a different integer (say, between 1 and 10) to let it know which thread it is.

After initializing a new Thread, we can start it running with the `Fork` method. This doesn't immediately begin the thread execution; instead it makes the thread `READY` to run and places it on the `readyList`. The `readyList` is a linked list of Threads. All Threads on the `readyList` have status `READY`. `ReadyList` is a global variable. There is another global variable named `currentThread`, which points to the currently executing Thread object; i.e., the Thread whose status is `RUNNING`.

The Yield method should only be invoked on the currently running thread. It will cause a switch to some other thread.

Follow the code in Yield closely to see what happens when a thread switch occurs. First, interrupts are disabled; we don't want any interference during a thread switch. The readyList and currentThread are shared variables and, while switching threads, we want to be able to access and update them safely. Then Yield will find the next thread from the readyList. (If there is no other thread, then Yield is effectively a nop.) Then Yield will make the currently running process READY (i.e., no longer RUNNING) and it will add the current thread to the tail end of the readyList. Finally, it will call the Run function to do the thread switch.

The Run method will check for stack overflow on the current thread. It will then call Switch to do the actual Switch.

Switch may be the most fascinating function you ever encounter! It is located in the assembly code file Switch.s, which you should look at carefully. Switch does not return to the function that called it. Instead, it switches to another thread. Then it returns. Therefore, the return happens to another function in another thread!

The only place Switch is called is from the Run function, so Switch returns to some invocation of the Run function in some other thread. That copy (i.e., invocation) of Run will then return to whoever called it. This could have been some other call to Yield, so we'll return to another Yield which will return to whoever called it.

And this is exactly the desired functionality of Yield! A call to Yield should give up the processor for a while, and eventually return after other threads have had a chance to execute.

Run is also called from Sleep, so we might be returning from a call to Sleep after a thread switch.

**Question 1:** How is everything set up when a thread is first created? How can we "return to a function" when we have not ever called it? Take a look at function ThreadStartMain in file Thread.k and look at function ThreadStartUp in file Switch.s.

**Question 2:** What happens when a thread is terminated? Take a look at ThreadFinish in file Thread.k. How is the thread put to sleep with no hope of ever being awakened? (No wonder they call it "Thread Death!")

Next, take a look at what happens when a Timer interrupt occurs while some thread is executing. This is an interrupt, so the CPU begins by interrupting the current routine's execution and pushing some state onto its (system) stack. Then it disables interrupts and jumps to the assembly code routine called TimerInterruptHandler in Runtime.s, which just calls the TimerInterruptHandler function in Thread.k.

In TimerInterruptHandler, we call Yield, which then switches to another thread. Later, we'll come back here, when this thread gets another chance to run. Then, we'll return to the assembly language routine which will execute a "reti" instruction. This will restore the state to exactly what it was before and the interrupted routine (whatever it was) will get to continue.

Note that this code maintains a variable called currentInterruptStatus. This is because it is rather difficult to query the "I" bit of the CPU. It is easier to just change the variable whenever a change to the interrupt status changes. We see this occurring in the TimerInterruptHandler function. Clearly interrupts will be disabled immediately after the interrupt occurs. And the Yield function will preserve the interrupt status. So when we return from Yield, interrupts will once again be disabled. Before returning to the interrupted thread, we set the currentInter-

ruptStatus to ENABLED. (They must have been enabled before the interrupt occurred-or else it could not have occurred-so after we execute the "reti" instruction, the status will revert to what it was before, namely ENABLED.)

Now you are ready to start playing with and modifying the code! Please experiment with the code we have just discussed, as necessary to understand it.

## 4 Run the "SimpleThreadExample" Code

Execute and trace through the output of SimpleThreadExample in file Main.k.

In TimerInterruptHandler there is a statement

```
printChar ('_')
```

which is commented out. Try uncommenting it. Make sure you understand the output.

**Question 3:** In TimerInterruptHandler, there is a call to Yield. Why is this there? Try commenting this statement? What happens. Make sure you understand how Yield works here.

## 5 Run the "MoreThreadExamples" Code

Trace through the output. Try changing this code to see what happens.

## 6 Implement the "Mutex" Class

In this part, you must implement the class Mutex. The class specification for Mutex is given to you in Synch.h:

```
class Mutex
  superclass Object
  methods
    Init ()
    Lock ()
    Unlock ()
    IsHeldByCurrentThread () returns bool
endClass
```

You will need to provide code for each of these methods. In Synch.k you'll see a behavior construct for Mutex. There are methods for Init, Lock, Unlock, and IsHeldByCurrentThread, but these have dummy bodies. You'll need to write the code for these four methods.

You will also need to add a couple of fields to the class specification of Mutex to implement the desired functionality.

How can you implement the Mutex class? Take a close look at the Semaphore class; your implementation of Mutex will be quite similar.

First consider the IsHeldByCurrentThread method, which may be invoked by any thread. The code of this method will need to know which thread is holding a lock on the mutex; then it can compare that to the currentThread to see if they are the same. So, you might consider

adding a field (perhaps called `heldBy`) to the `Mutex` class, which will be a pointer to the thread holding the mutex. Of course, you'll need to set it to the current thread whenever the mutex is locked. You might use a null value in this field to indicate that no thread is holding a lock on the mutex.

When a lock is requested on the mutex, you'll need to see if any thread already has a lock on this mutex. If so, you'll need to put the current process to sleep. For putting a thread to sleep, take a look at the method `Semaphore.Down`. At any one time, there may be zero, one, or many threads waiting to acquire a lock on the mutex; you'll need to keep a list of these threads so that when an `Unlock` is executed, you can wake up one of them. As in the case of `Semaphores`, you should use a FIFO queue, waking up the thread that has been waiting longest.

When a mutex lock is released (in the `Unlock` method), you'll need to see if there are any threads waiting to acquire a lock on the mutex. You can choose one and move it back onto the `readyList`. Now the waiting thread will begin running when it gets a turn. The code in `Semaphore.Up` does something similar.

It is also a good idea to add an error check in the `Lock` method to make sure that the current thread asking to lock the mutex doesn't already hold a lock on the mutex. If it does, you can simply invoke `FatalError`. (This would probably indicate a logic error in the code using the mutex. It would lead to a deadlock, with a thread frozen forever, waiting for itself to release the lock.) Likewise, you should also add a check in `Unlock` to make sure the current thread really does hold the lock and call `FatalError` if not. You'll be using your `Mutex` class later, so these checks will help your debugging in later assignments.

The function `TestMutex` in `Main.k` is provided to exercise your implementation of `Mutex`. It creates 7 threads that compete vigorously for a single mutex lock.

## 7 Implement the Producer-Consumer Solution

Your textbook, Silberschatz, et al. contains a discussion of the Producer-Consumer problem, also known as the Bounded-Buffer Problem, including a solution. Implement this in KPL using the classes `Mutex` and `Semaphore`. Deal with multiple producers and multiple consumers, all sharing a single bounded buffer.

The `Main` package contains some code that will serve as a framework. The buffer is called `buffer` and contains up to `BUFFER_SIZE` (e.g., 5) characters. There are 5 producer processes, each modeled by a thread, and 3 consumer processes, each modeled by a thread. Thus, there are 8 threads in addition to the main thread that creates the others.

Each producer will loop, adding 5 characters to the buffer. The first producer will add five 'A' characters, the second producer will add five 'B's, etc. However, since the execution of these threads will be interleaved, the characters will be added in a somewhat random order.

## 8 Implement the Dining Philosopher's Solution Using a Monitor

A starting framework for your solution is provided in `Main.k`. Each philosopher is modeled with a thread and the code we've provided sets up these threads. The synchronization will be



controlled by a "monitor" called ForkMonitor.

The code for each thread/philosopher is provided for you. Look over the PhilosophizeAndEat method; you should not need to change this code.

The monitor to control synchronization between the threads is implemented with a class called ForkMonitor. The following class specification of ForkMonitor is provided:

```
class ForkMonitor
  superclass Object
  fields
    status: array [5] of int      -- For each philosopher: HUNGRY,
                                -- EATING, or THINKING

  methods
    Init ()
    PickupForks (p: int)
    PutDownForks (p: int)
    PrintAllStatus ()
endClass
```

You'll need to provide the code for the Init, PickupForks and PutDownForks methods. You'll also need to add additional fields and perhaps even add another method.

The code for PrintAllStatus is provided. You should call this method whenever you change the status of any philosopher. This method will print a line of output, so you can see what is happening.

How can you proceed? You'll need a mutex to protect the monitor itself. There are two main methods (PickupForks and PutDownForks) which are called by the philosopher threads. Upon beginning each of these methods, the first thing is to lock the monitor mutex. This will ensure that only one thread at a time is executing within the monitor. Just before each of these methods returns, it must unlock the monitor (by unlocking the monitor's mutex) so that other threads can enter the monitor code.

You'll also need to use the Condition class, which is provided in the Synch package. (The Condition class uses the class Mutex, so it is assumed that you've finished and tested the Mutex class.)

The BLITZ emulator has a number of parameters and one of these is how often a timer interrupt occurs. The default value is every 5000 instructions. You might try changing this parameter to see how it affects your programs behavior. To change the simulation parameters, type the *sim* command into the emulator. This command will give you the option to create a file called

.blitzrc

After creating this file, you can edit it by hand. The next time you run the emulator, it will use this new value. Also note that too small a value-like 1000-will cause the program to hang.

**Question 4:** What do you suppose causes this effect?

**Question 5:** Find and read about "Hoare semantics" and "Mesa-style semantics". The comments in the code in Synch.k say that this version implements "Mesa-style" semantics. Is this the same or different from Hoare semantics? If different, what is the difference? Quote the code that implements the "Mesa-style semantics."

## 9 An Example of Correct Output

The following files contain an example of what correct output should look like:

```
DesiredOutput1.pdf
DesiredOutput2.pdf
DesiredOutput3.pdf
DesiredOutputGame.pdf (for the next section)
```

## 10 Implement the Gaming Parlor Problem (50 points)

Here is the problem scenario: groups of customers come in to a "gaming parlor" to play games. They go to the front desk to obtain one or more dice, which are used by the group while they are playing their game, and then returned to the front desk. The front desk is in charge of lending out the dice and collecting them after each game is finished.

The gaming parlor owns only 8 dice, which are available at the front desk before the first group comes in.

The customers can play the following games. Listed after each game in parentheses is the number of dice required to play that game.

```
Backgammon (4)
Risk (5)
Monopoly (2)
Pictionary (1)
```

You should model the front desk as a monitor with the following entry methods:

```
Request (numberOfDice: int)
Return (numberOfDice: int)
```

Model each group of customers as a thread. When a group is ready to play, it must obtain the necessary number of dice. If the required number of dice is not available, then the group (i.e., the thread) must wait. You might use a condition variable that "more dice have become available."

You should model the following eight different groups. Each group plays one game, as shown below, but each group plays its game 5 times. Each group must return their dice after each game and then re-acquire the dice before playing again.

```
A - Backgammon
B - Backgammon
C - Risk
D - Risk
E - Monopoly
F - Monopoly
G - Pictionary
H - Pictionary
```

To simulate playing the game, simply call the Yield method.

This problem is a generalization of the problem of resource allocation where (1) there are a number of resources (dice) but each is identical; (2) every requesting process needs a one or more units of the resource, (3) each requesting thread knows how many units it will need before requesting any units and that info is included in the request, (4) all units are returned before any further requests are made.

In a monitor, each method is either an "entry" method or a local method. A monitor will always have one Mutex lock, which is associated with entering the monitor. Every monitor will also have zero or more condition variables, as required by the particular application.

An entry method must always begin by locking the monitor's mutex. Before it returns, it must always unlock the mutex. Within the entry method, the code may Signal some of the condition variables and may Wait on condition variables. (In some applications, it may make sense to Broadcast to a condition variable.) However, once inside of an entry method, the code should never touch the monitor's mutex lock, or try to look inside the condition variables.

You should model the front desk as a monitor and you should use condition variables instead of semaphores.

If deadlock occurs, the program will freeze up and some requests for dice will go unsatisfied forever. This is a disaster! Regardless of the order in which the groups make their requests, your solution should be structured such that deadlock can never occur.

Your solution must not be subject to any race conditions. In other words, regardless of the order in which the groups make their requests and return their dice, each die must never be allocated to more than one group at a time. It should never be the case that groups are allowed to proceed when there are too few dice. Likewise, if a group has returned its dice, other groups which are waiting must be allowed to proceed once enough dice have become available.

Starvation can occur if it is possible that one thread can be delayed infinitely by other threads' requests. If the game parlor problem is extended to assume that each group will continue to play game after game forever, then starvation might be possible, if you are not careful in your solution. If some group X comes in requesting a lot of dice, it will be made to wait until enough dice are available. If it is possible that other groups can come in, request a small number of dice, and have their requests granted, then group X might get delayed forever, since there are never enough dice at the front desk at once to satisfy group X's needs. In other words, it might be possible for starvation of X to occur. In your solution starvation should not be possible. (Of course, with each group limited to playing only 5 games, all the outstanding dice will always get returned eventually and starvation can never occur, but your solution should also avoid starvation in the presence of a never-ending sequence of Request and Return operations.)

To verify that your code is working, please insert print statements to produce output like this...

```
Initializing Thread Scheduler...
A requests 4
-----Number of dice now avail = 8
A proceeds with 4
-----Number of dice now avail = 4
```

```

B requests 4
-----Number of dice now avail = 4
B proceeds with 4
-----Number of dice now avail = 0
D requests 5
-----Number of dice now avail = 0
E requests 2
-----Number of dice now avail = 0
A releases and adds back 4
-----Number of dice now avail = 4
B releases and adds back 4
-----Number of dice now avail = 8
C requests 5
-----Number of dice now avail = 8
H requests 1
-----Number of dice now avail = 8
B requests 4
-----Number of dice now avail = 8
D proceeds with 5
-----Number of dice now avail = 3
...etc...

```

This output makes it fairly easy to see what the program is doing and verify that it is correct.

Below is a method you should use to produce your output. Feel free to copy this method straight into your program.

```
behavior GameParlor
```

```
... Other methods...
```

```

method Print (str: String, count: int)
--
-- This method prints the current thread's name and the arguments.
-- It also prints the current number of dice available.
--
    print (currentThread.name)
    print (" ")
    print (str)
    print (" ")
    printInt (count)
    nl ()
    print ("-----Number of dice now avail = ")
    printInt (numberDiceAvail)

```

```
    nl ()
endMethod
```

```
endBehavior
```

This method would be called in several places...

- At the beginning of the Request method: self.Print ("requests", numNeeded)
- At the end of the Request method: self.Print ("proceeds with", numNeeded)
- In the Return method: self.Print ("releases and adds back", numReturned)

## 11 Implement the Sleeping Barber Problem (50 points)

An older edition of the Tanenbaum textbook describes the “Sleeping Barber” problem and gives a solution in “C”. (This material can also be found in a separate file called SleepingBarber-Problem.pdf in the a2 directory on the class web site.)

Translate this into a working KPL program. You’ll also need to create some code to print out what happens when you run the program.

What will you do for the “cut\_hair” and the “get\_haircut” methods? You’ll need at least one print statement in each routine, but be careful: they both should run at more-or-less the same time.

One barber is okay, but how many customers will you model? Will you just throw a bunch of customers at the barbershop all at once? This might not test your code very well. You might need to try several runs where customers come in at different intervals. For example, if there are enough customers so the barber is always busy, you never test the code for the barber going to sleep. If there is only one customer at a time and the barber always sleeps between customers, you may not test cutting two customers in a row. And finally, if you don’t have customers arriving fast enough, you may not test having the waiting room full. How can you test all these cases?

How long will the haircut last? You can implement waiting with a busy loop, such as

```
for i = 1 to 100
  .. want to yield here?...
endFor
```

## 12 What to Hand In

Complete all the programming requested above. In a text file or word processed document, answer all five questions asked in the above text.

You may put all your sources for this assignment in an “a2” branch in your git repository. Put the files in an “a2” directory. This “a2” directory should not be on the master branch. You need at minimum the following files in your “a2” directory.

`Synch.h`, `Synch.k`, `Main.k`, `Game.*`, `Barber.*`

Also capture output for sections 6, 7, 8, 10 and 11 in separate files named **section-X** where X is the section number. Add and commit those files to your “a2” directory. Remember to push all your changes to gitlab. If you want comments on your code rather than just notations on your cover page of points lost, print the code you wrote. You may have to print some other code, but a good yellow highlighter can easily show the code you wrote. A highlighted box around the code is good. You don’t need to highlight every line of code. `a2ps(1)` produces great output that is very readable.

Finally, print the answers to the 5 questions in this writeup and turn them in. Also, provide a convincing argument that your Barbers implementation is correct. Please do not turn in handwritten answers. **The answers must not be comments in your code.**

Remember to have a cover page with your name, this class and the assignment number.

## 13 Basis for Grading

The code you submit will be graded on both correctness and style. Correctness means that the code must work right and not contain bugs. Style means that the code must be neat, well commented, well organized, and easy to read. In style, your code should match the code we are distributing to you. The indentation should be the same and the degree of commenting should be the same.